中国科学院
CHINESE ACADEMY OF SCIENCES

SciSpace

Towards a Serverless-Style Cloud Service for Scientific Data Lifecycle Management

(Part of the project SCIBDS)

Huajin Wang, Zhihong Shen, Yuepeng Li
wanghj@cnic.cn

Nov. 30, 2018
BigSDM 2018, Beijing

中国科学院
计算机网络信息中心
Computer Network Information Center,
Chinese Academy of Sciences

Outline

- Data Lifecycle Management
  - What is
  - Traditional Approach
  - Serverless Approach
- SciSpace: Our Serverless Solution
  - Space: The Unified Serverless UI
  - Packone (Part I): Make Engine Elastic
  - Packone (Part II): Make URI Highly Available
- Conclusion

What is The Data Lifecycle Management

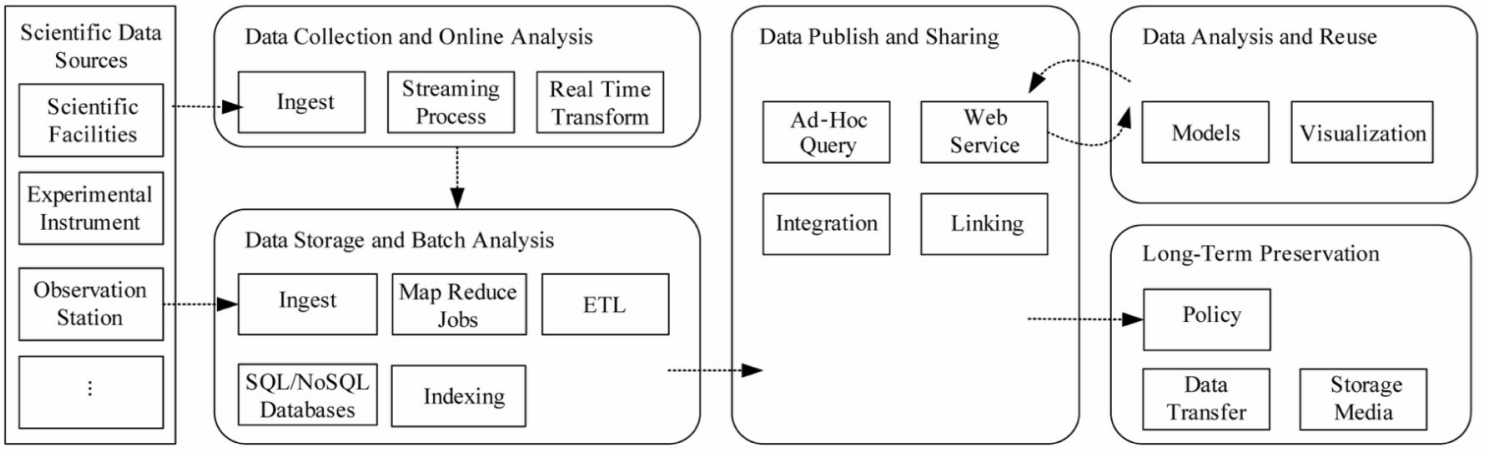

Rethink:

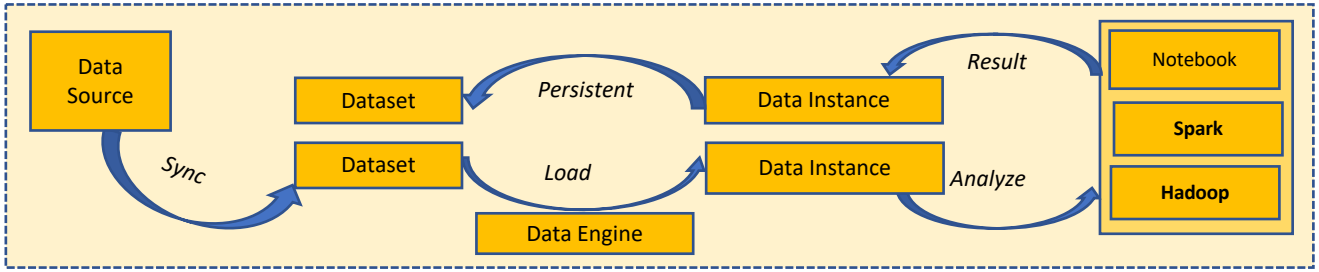

Data Lifecycle Management: The Traditional Approach

- Data sources and scientists are scattered in different locations/clouds
  - data analysis always need to access remote data stores
  - data trans between Heterogenous Infrastructure are complicated
- No unified data lifecycle management environment
  - Scientists need to using different level UI to access data or computing powers
    - SSH/FTP/HDFS/S3/
  - Dataset or analysis results can hardly be referenceable by peers

Data Lifecycle Management: The Serverless Approach

What

- A unified data lifecycle management environment
  - Data Sync/Load/Share
  - Data Query/Analysis
- Server location is insignificant
- Server login is not required

Advantages

- Liberate Scientists

Challenges

- The lifecycle challenge
  - How to takeover data lifecycle management
    - without direct interactions with underlying servers
- The scalability challenge
  - How to automatically scale computing infrastructure
    - to meet the performance constraints
- The HA challenge
  - How to tolerant server failures
    - to keep service highly available

SciSpace: Our Serverless Data Lifecycle Management Service

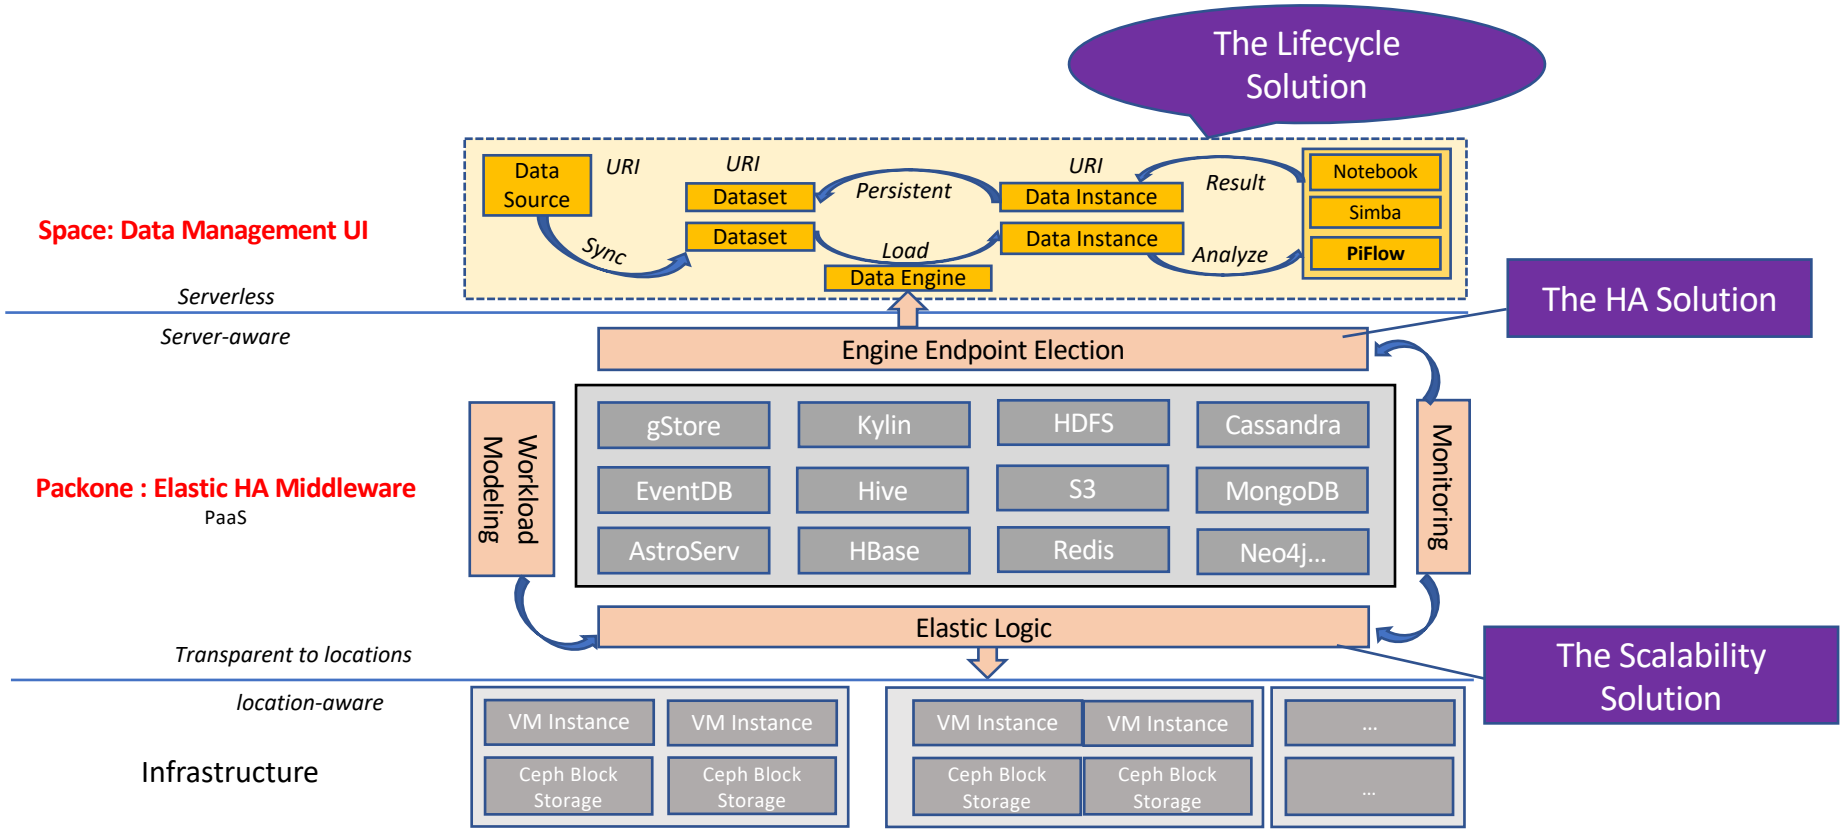

Space: The Unified Serverless UI

- Easy to share dataset
- Dataset can be loaded into data instances multiple times
- Using the URI field to talk with each other.
- Monitor underlying computing resource located anywhere

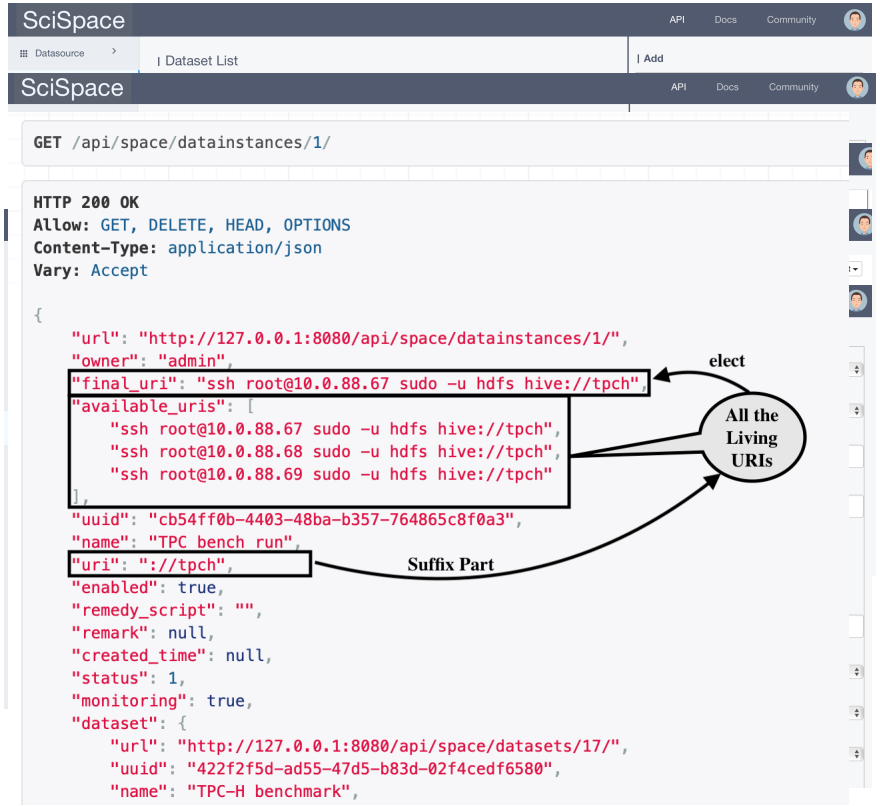

Packone: Make Engine Elastic

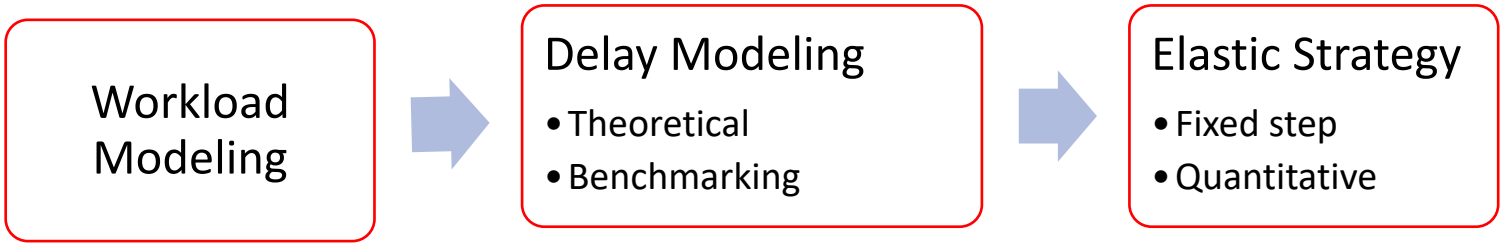

Packone: Make Engine Elastic

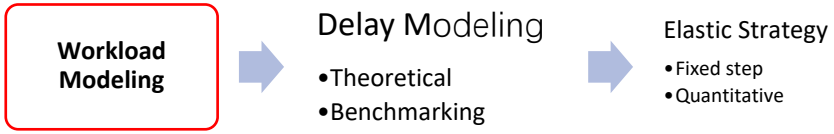

- Survey Approach
  - Portable but easy to be retired
- Generation Approach
  - Up to date but engine dependent
- Our Evolutionary Approach
  - Portable
  - Up to date
  - Engine independent

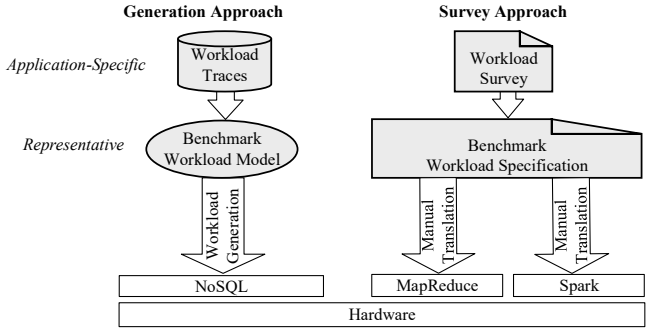

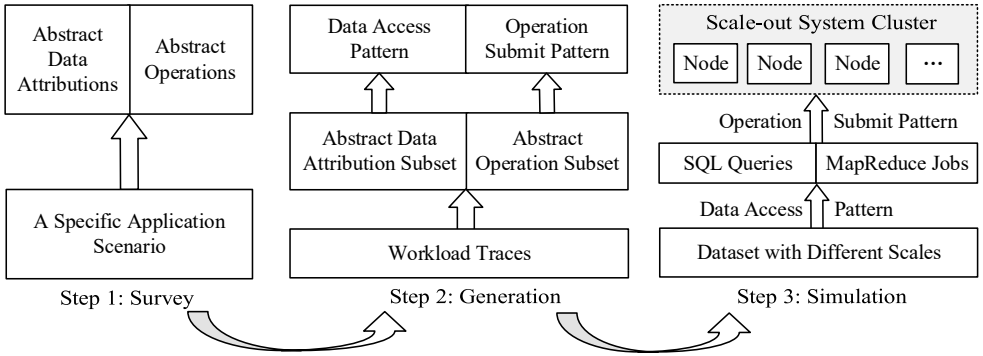

Packone: Make Engine Elastic

Workload Modeling → **Delay Modeling** (•**Theoretical** •**Benchmarking**) → Elastic Strategy (•Fixed step •Quantitative)

- **Theoretical**: LT approx. depicted how node qty. replication factor & consistency level affect the operation delay

Theorem (The LT approximation of the expected sojourn time of the (n, k) fork-join queues)

$$\underbrace{\overbrace{t_{(n,k)}}^{(n,k)\ queue}}_{order\ statistic} = \sum_{i=k}^{n} W_i^{n,k} \underbrace{\overbrace{t_{(i,i)}}^{(i,i)\ queue}}_{maximum}$$

Proof.

$t_{1..n}$ *are jointly-identically distributed* $\Longrightarrow t_{(n,k)} = \sum_{i=k}^{n} W_i^{n,k} t_{(i,i)}$. □

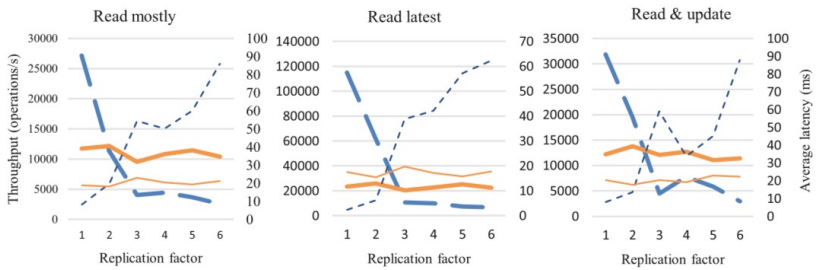

- **Benchmarking** gives the real delay of the target engine
- **Real-life delay** ≈ {**Benchmarking@Workload Model**} × **Theoretical**

Packone: Make Engine Elastic

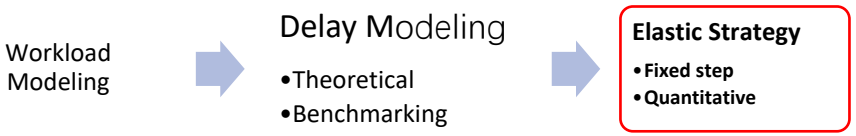

- Elastic Strategies to keep a satisfactory delay
  - Fixed Step: Easy to employ but low-precision
  - Quantitative: Fast-forward but hard to calculate
  - Our Mixed Strategy
    - Beginning: Fixed Step
    - Ending: Quantitative

Theorem (The LT approximation of the expected sojourn time of the (n, k) fork-join queues)

$$\overbrace{\underbrace{t_{(n,k)}}_{\text{order statistic}}}^{(n,k)\ \text{queue}} = \sum_{i=k}^{n} W_i^{n,k} \overbrace{\underbrace{t_{(i,i)}}_{\text{maximum}}}^{(i,i)\ \text{queue}}$$

Proof.

$t_{1..n}$ *are jointly-identically distributed* $\Longrightarrow t_{(n,k)} = \sum_{i=k}^{n} W_i^{n,k} t_{(i,i)}$. □

Packone: Make URI Highly Available

- Restful API
- Realtime available URI set
- URI Election Strategy
  - Naïve: Pick out the 1st living one
  - Load balance (TODO)
  - Round Robin (TODO)

```
GET /api/space/datainstances/1/

HTTP 200 OK
Allow: GET, DELETE, HEAD, OPTIONS
Content-Type: application/json
Vary: Accept

{
    "url": "http://127.0.0.1:8080/api/space/datainstances/1/",
    "owner": "admin",
    "final_uri": "ssh root@10.0.88.67 sudo -u hdfs hive://tpch",
    "available_uris": [
        "ssh root@10.0.88.67 sudo -u hdfs hive://tpch",
        "ssh root@10.0.88.68 sudo -u hdfs hive://tpch",
        "ssh root@10.0.88.69 sudo -u hdfs hive://tpch"
    ],
    "uuid": "cb54ff0b-4403-48ba-b357-764865c8f0a3",
    "name": "TPC bench run",
    "uri": "://tpch",
    "enabled": true,
    "remedy_script": "",
    "remark": null,
    "created_time": null,
    "status": 1,
    "monitoring": true,
    "dataset": {
        "url": "http://127.0.0.1:8080/api/space/datasets/17/",
        "uuid": "422f2f5d-ad55-47d5-b83d-02f4cedf6580",
        "name": "TPC-H benchmark",
```

Conclusion

- Traditional approach
  - Scattered dataset and computing resources
  - Complicated and difficult
  - Poorly referenceable
- Serverless Approach
  - The lifecycle challenge
  - The scalability challenge
  - The HA challenge
- SciSpace: Our Serverless Data Lifecycle Management Service
  - Space: The Unified Serverless UI
  - Packone Make Engine Elastic and URI Highly Available

References

- J Li , Z Shen, and X Meng. Scientific Big Data Management: Concepts, Technologies and System.
- H Wang, J Li, H Zhang, Y Zhou. Benchmarking replication and consistency strategies in cloud serving databases: Hbase and cassandra
- H Wang, J Li, Z Shen, Y Zhou. Approximations and Bounds for (n, k) Fork-Join Queues: A Linear Transformation Approach. CCGrid 2018.
- H Wang, M Wan, R Han etc. Towards the Automatic Evolution of Workload Models in Large-scale Astronomical Data Management.

Thanks!

Q & A